SHIP BUILDER PROOF

AI agents are moving faster than approval systems.

The risk is not that agents can act. The risk is that nobody can show who approved the movement, what proof changed, and what stayed blocked.

RAW AGENT MOTION

Tool request
comment, create label, repo admin

Policy gap
intent is not evidence

approval receipt missing

DECISION INTEGRITY

GitHub hashes code. IA hashes the decision to move.

Code state is necessary. It is not the same as approval state.

CODE STATE

commit hash

a4c4e94

what changed in code

DECISION STATE

packet hash

sha256:8fb9...c71e

why movement was allowed, reviewed, or blocked before anything moved

CATEGORY WEDGE

GitHub proves what shipped. IA proves why it was safe to move.

ONE REVIEWRUN

One ReviewRun turns raw intent into a packet.

Raw intent is not trusted. Proof changes packet state. Downstream systems consume the packet.

Agent intent

support-triage-bot asks for repo access

->

IA Packet

packet id + revision + decision hash

->

Movement verdict

allowed, review required, blocked

SCOPED OUTCOME

read issues -> allowed   comment -> review   admin -> blocked

PORTABLE RECEIPT

Humans approve scope. IA records the receipt.

CTO, Security, and Procurement approve the human scope. IA does not approve by chat; it records the packet-backed receipt.

APPROVAL RECEIPT

ia-review-run-receipt-7619782b

sha256:6c40e9f7b1d4...83ac

Allowed: selected repo only

CTO scope verified

Security verified

Procurement verified

DECISION BOUNDARY

Still blocked: repo admin, org-wide write, secrets.

DOWNSTREAM GATE

Portkey consumes the packet before movement.

IA is the packet authority. Portkey is the enforcement gateway. No admin mutation is required for the demo proof.

IA Packet
packet id, revision, allowed scope

->

BYO Guardrail
webhook asks IA for verdict

->

Portkey
enforces packet-backed state

rev_1 -> Block

rev_2 -> Allow with policy

API mutation: false

policy mutation: false

DEMO CLAIM

Portkey asks IA whether movement can proceed; IA returns the packet-backed verdict.

PROOFGRAPH

ProofGraph shows why the decision changed.

The graph is generated from the ReviewRun, not a screenshot. It carries packet revision, proof path, hash, and downstream state.

GRAPH HASH

8fb9...c71e

NEW CATEGORY

Portable approval receipts for AI movement.

InferenceAtlas lets a human approve scope once, then every downstream system verifies the same packet before movement.

SAFETY INVARIANTS

- IA approval granted: false
- GitHub write: false
- Portkey policy mutation: false

FINAL RECEIPT

safe to move

packet: rev_2

Allowed with policy. Still blocked stays explicit.

DEMO SENTENCE

Connect repo -> generate packet -> attach proof -> rerun -> Portkey consumes -> ProofGraph exports.